

# Best Practices in Machine Learning

CAS Artificial Intelligence - Machine Learning

**Tobias Mérinat**

Prof. Dr. Marc Pouly

[tobias.merinat@hslu.ch](mailto:tobias.merinat@hslu.ch)

# Motivation

1. You carefully split your data into training, validation and test set
2. You train your favorite learning algorithm on the training set
3. You perform hyperparameter optimization on the validation set
4. You get 20% error on the test set which seems rather high

What can you do ?

## Some Ideas ...

- You could try getting more or better training data
- You could try a smaller set of features
- You could try a larger set of features
- You could try entirely different features
- You could try more iterations of gradient descent or adapt the learning rate
- You could try other optimization algorithms than gradient descent
- You could try different hyper-parameters (e.g. regularization parameters)
- You could try other machine learning algorithms

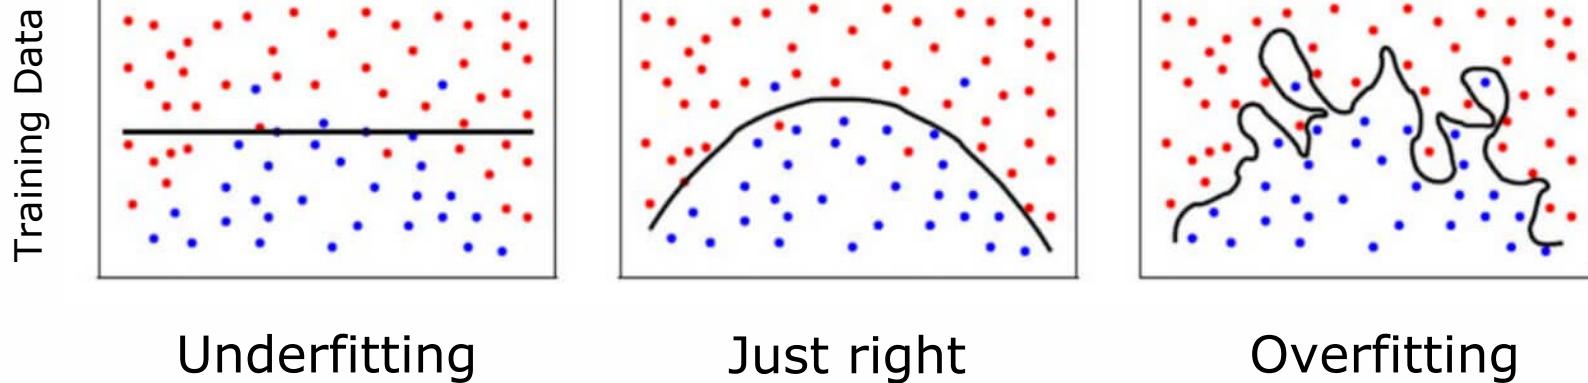
Just trying random modifications is not the way to go!  
You need to know what is going on first.

# Bias and Variance

Different sources of error can prevent supervised learning algorithms from generalizing beyond their training set

- The **bias** results from false assumptions in the learning algorithm. High bias can cause an algorithm to miss the relevant structure between features and target
- The **variance** results from sensitivity to small fluctuations in the training set. High variance can cause an algorithm to model the random noise in the training data rather than the intended outputs
- A model with high bias and low variance **underfits** the data
- A model with low bias and high variance **overfits** the data

# Underfitting vs Overfitting for Classification



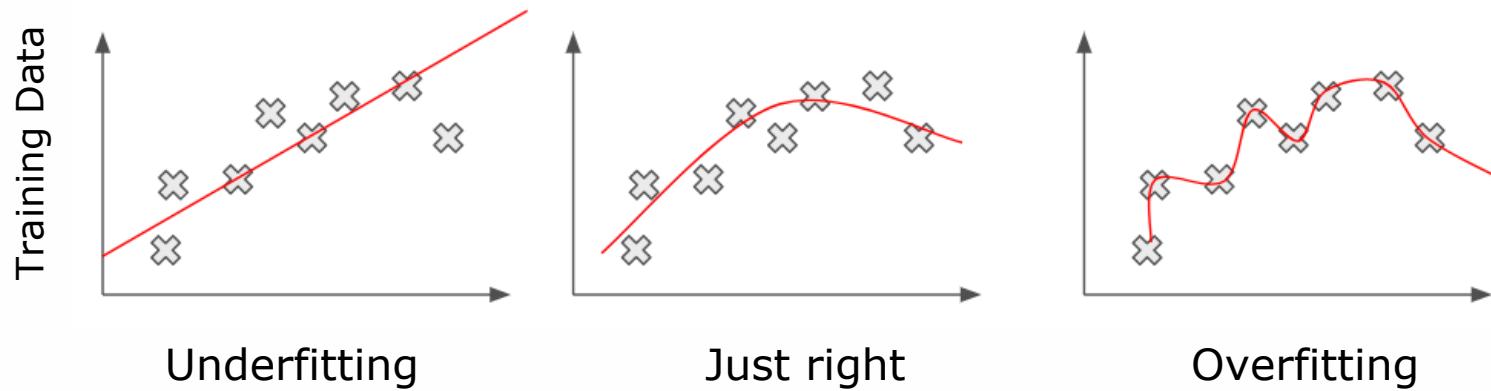
Underfitting:

- A straight line is just not complex enough to separate the classes
- This model performs badly on both training and test data

Overfitting:

- This model optimizes too much with respect to the training data
- It shows high performance on training but low performance on test data

# Underfitting vs Overfitting for Regression



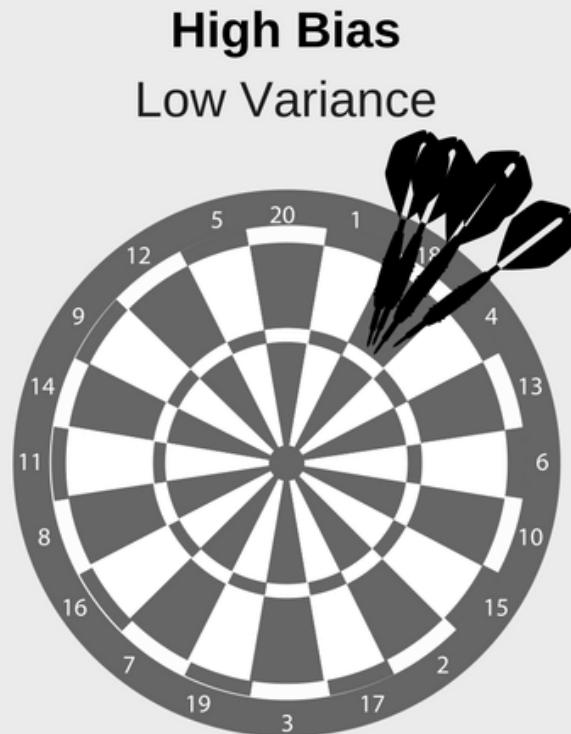
Underfitting:

- A straight line is just not complex enough to approximate the data
- This model performs badly on both training and test data

Overfitting:

- This model optimizes too much with respect to the training data
- It shows high performance on training but low performance on test data

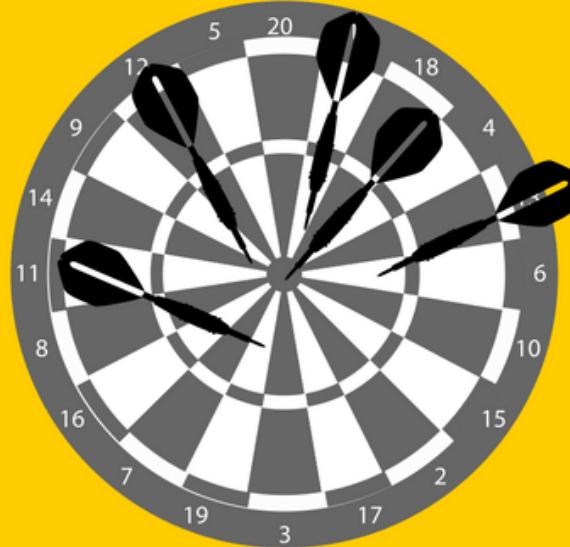
## Underfitting



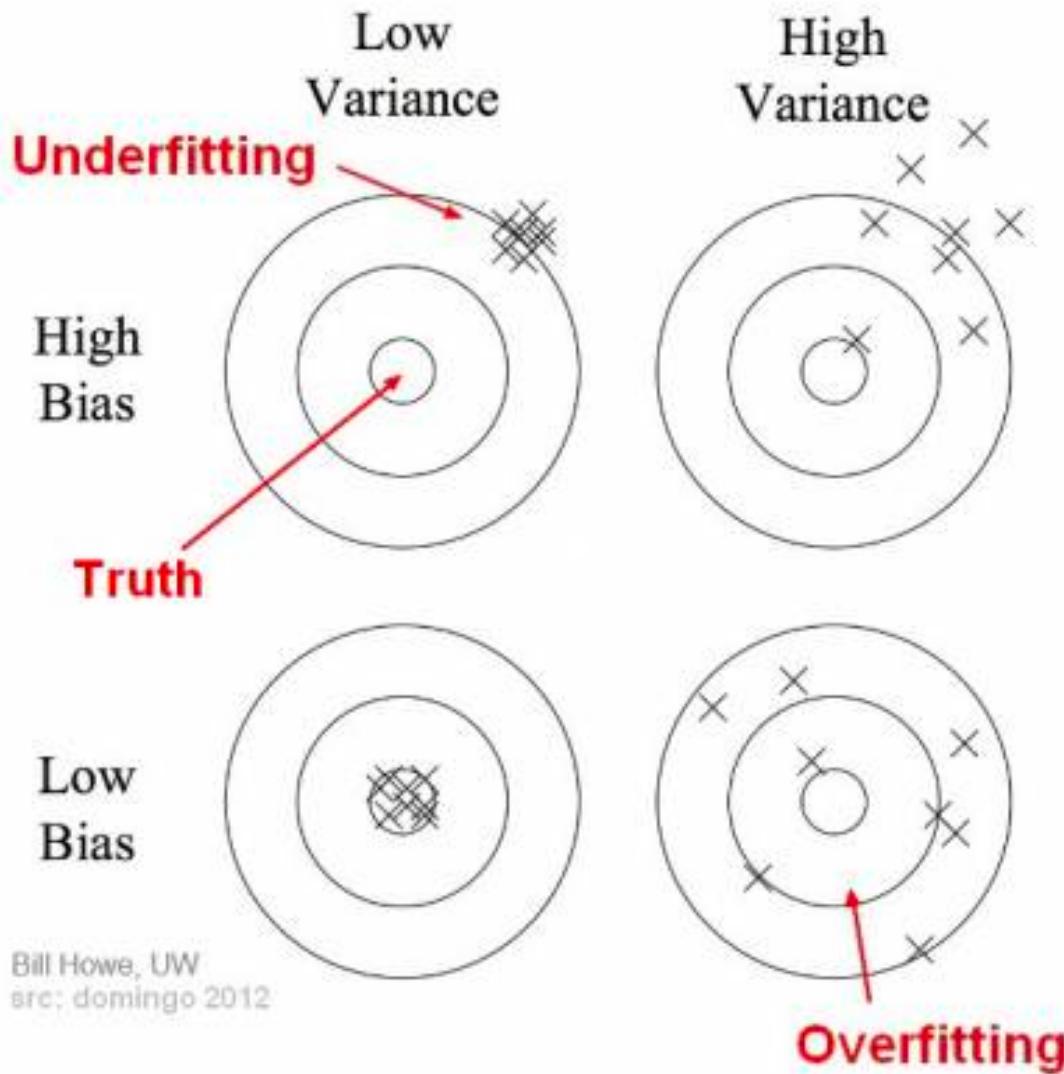
**High bias**, low variance  
algorithms train models that  
are consistent, but inaccurate  
*on average*.

## Overfitting

High Variance  
Low Bias

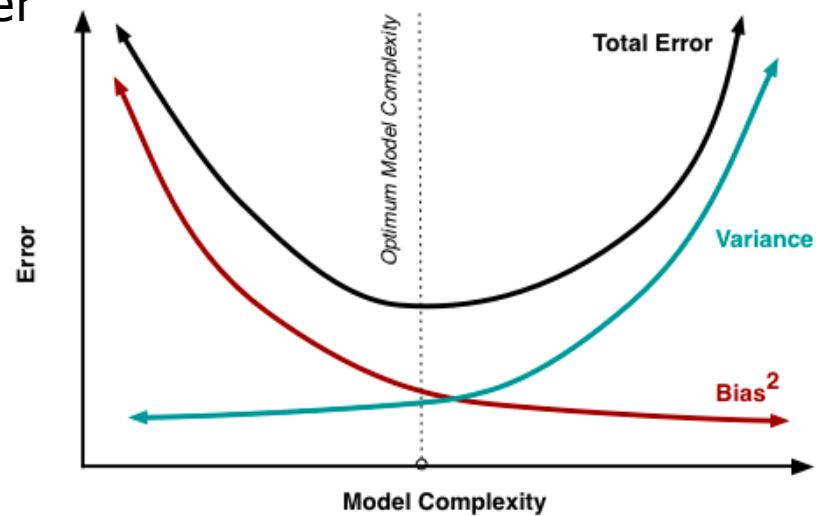


**High variance**, low bias  
algorithms train models that  
are accurate *on average*, but  
inconsistent.



# Bias-Variance Tradeoff

- Bias and Variance go hand-in-hand. It is often easy to improve one but simultaneously worsen the other
- This is called the **Bias-Variance Tradeoff**
- But our goal is to reach an overall improvement!



*Note that the Bias-Variance Tradeoff is less severe with large deep neural networks. There, simply adding more and more and more data often improves variance while not affecting bias negatively.*

Image source: [dataquest.io/](https://dataquest.io/)

## Our List of Ideas revisited...

- You could try getting more or better training data
- You could try a smaller set of features
- You could try a larger set of features
- You could try more iterations of gradient descent or adapt the learning rate
- You could try other optimization algorithms than gradient descent
- You could try different hyper-parameters (e.g. regularization parameters)
- You could try other machine learning algorithms
- 

Which action would be good when underfitting?  
When overfitting?

# Recap: More Complex Evaluation Workflows



- A recommended split is 60% / 20% / 20%
- How to evaluate a machine learning classifier properly:
  1. Loop over all hyperparameter combinations you would like to try
  2. Train model on selected hyperparameter setting and training set
  3. Evaluate model on validation set and measure performance
  4. Select model with best performance as candidate
  5. Evaluate model on test set for final performance estimate
- This workflow requires a lot of data

# Recap: Cross Validation

A train-validation-test split of 60% / 20% / 20% is common\*, but this only works when there is enough data. If not, split your data into 80% training and 20% test data (lock your test data away) and perform **k-fold cross-validation**. Even in the presence of enough data, cross-validations may bring in additional trustworthiness

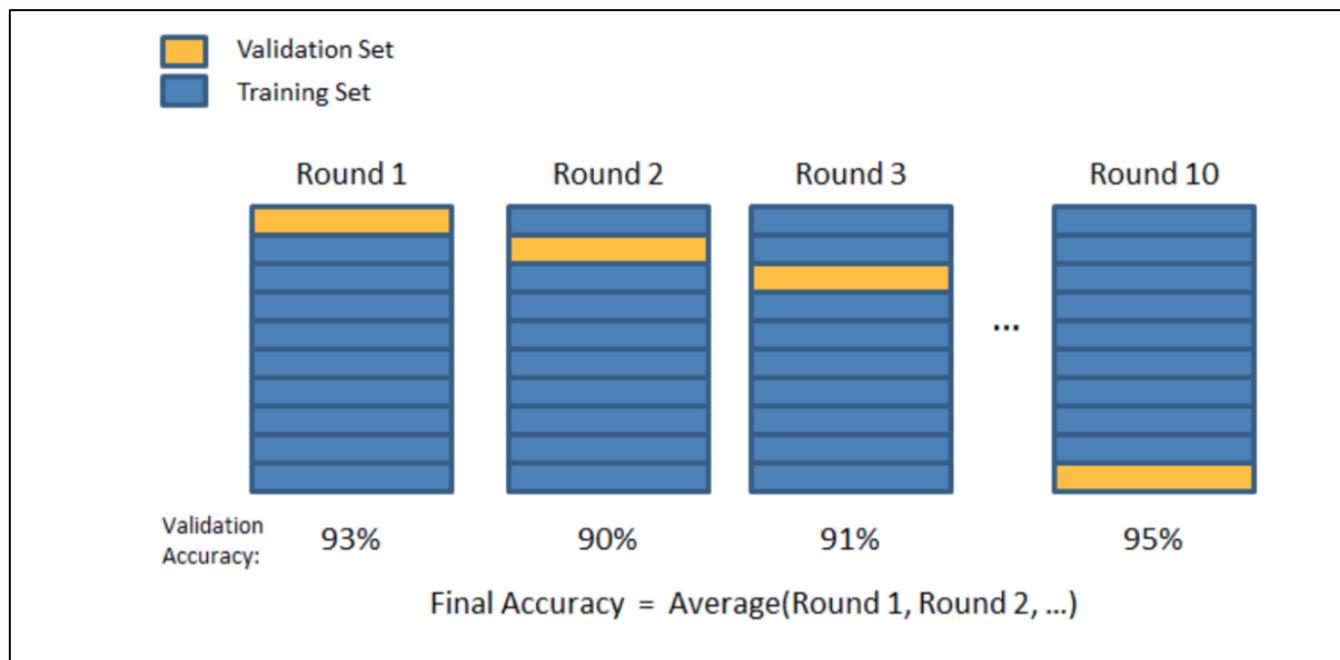
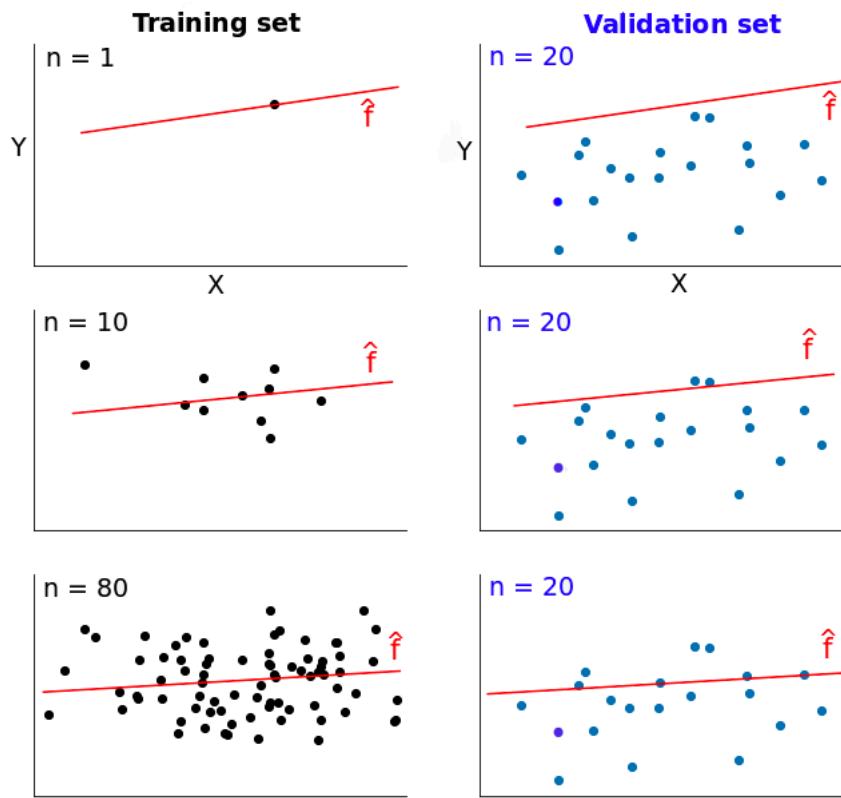


Illustration of 10-fold cross-validation

\*If you have a huge amount of data, as you should when training large deep NN, it is OK to make your validation and test sets a lot smaller than 20%.

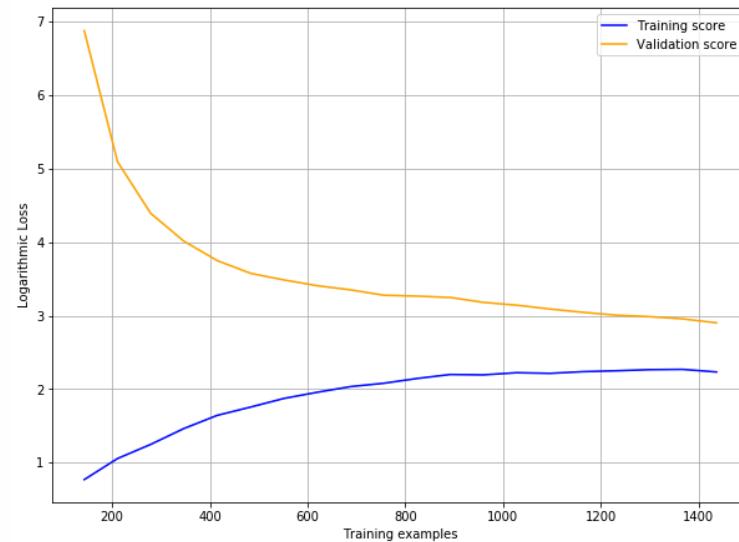
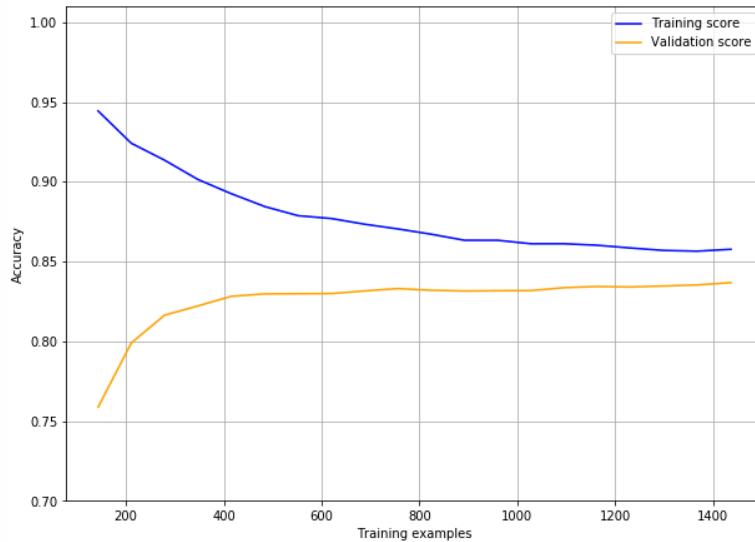
# Error in respect to number of samples



- Training error increases first and gets better over time
- Validation error is off in the beginning but approaches training error over time

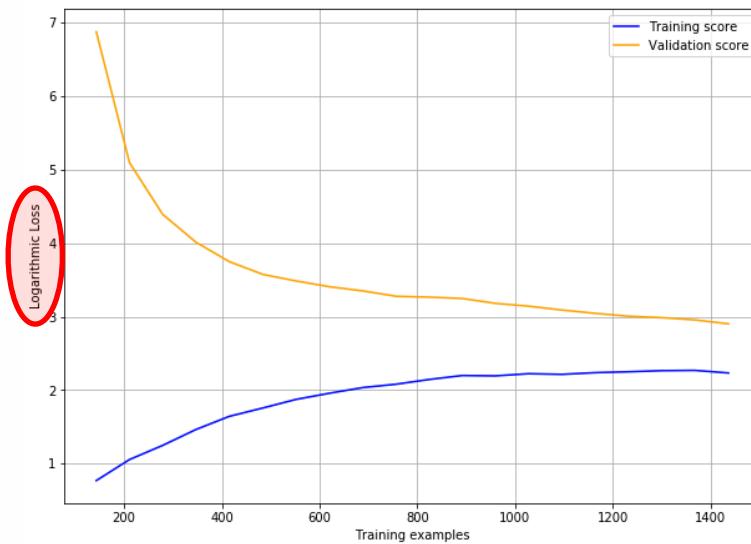
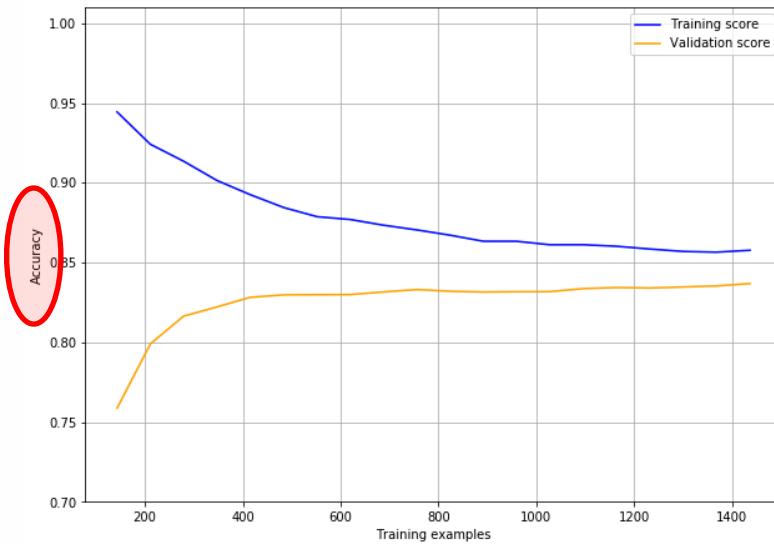
# Learning Curves

- Learning curves are a diagnostic tool
- The **x-axis** normally shows either «experience» or sometimes «model complexity»
- They usually show the «amount of learning» (in a chosen metric) on the **y-axis**



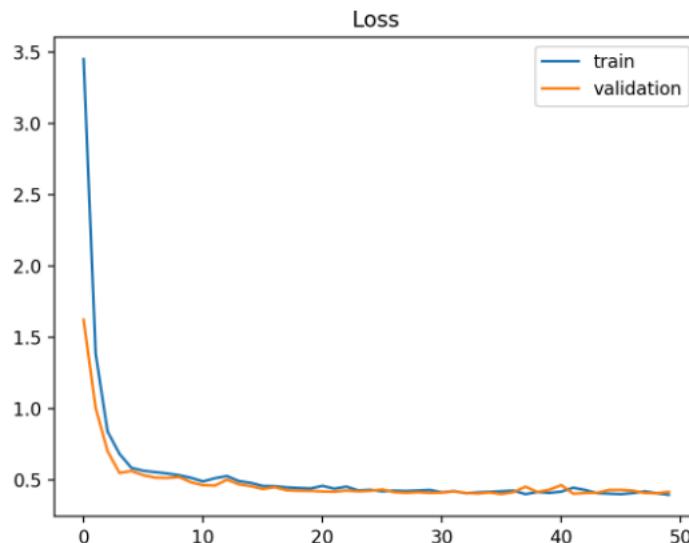
# Learning Curves

- Learning curves are a diagnostic tool
- The **x-axis** normally shows either «experience» or sometimes «model complexity»
- They usually show the «amount of learning» (in a chosen metric) on the **y-axis**



# Learning Curves in Neural Networks

- With Neural Networks, we often plot one data point per epoch (where above, we plotted one data point per batch)
- The x-axis still shows «experience», as we train longer (many epochs)
- But the effect of an increasing training error is much less pronounced, if it is there at all
- This makes the curves a bit easier to understand



A rather nice-looking set of learning curves

## Irreducible error

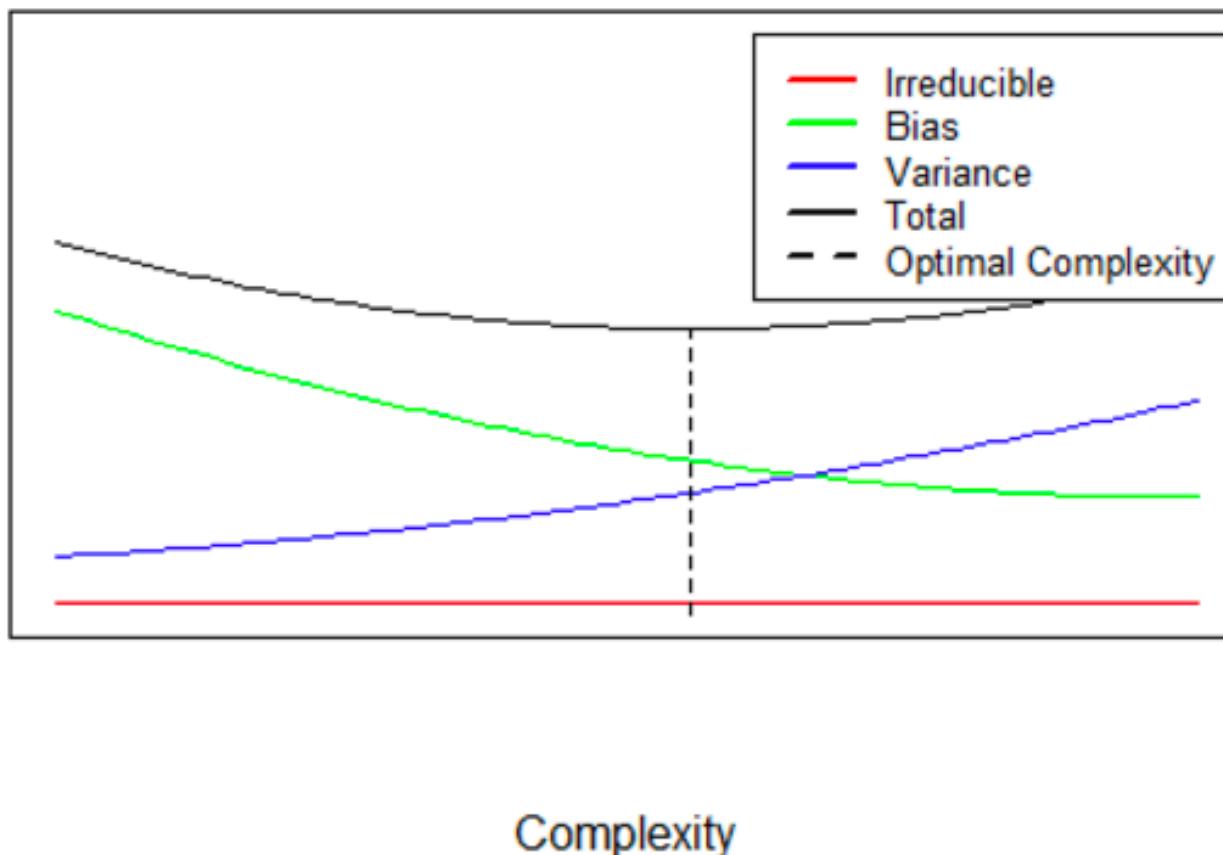
- With our two curves, do we have enough information to judge the situation?
- Not if the achieved training error is far from 0.
- There is a third source of error next to bias and variance:

The **irreducible error**       $Y = f(X) + \text{irreducible error}$

- It stems from noise in the data or lack of information (things that influence the outcome but are not represented in your data).
- We try to minimize the reducible error by lowering bias and variance, but we should also have a notion of what the irreducible error is, so we know what is achievable.

## Bias/Variance/Irreducible Error Decomposition

Contribution to Error

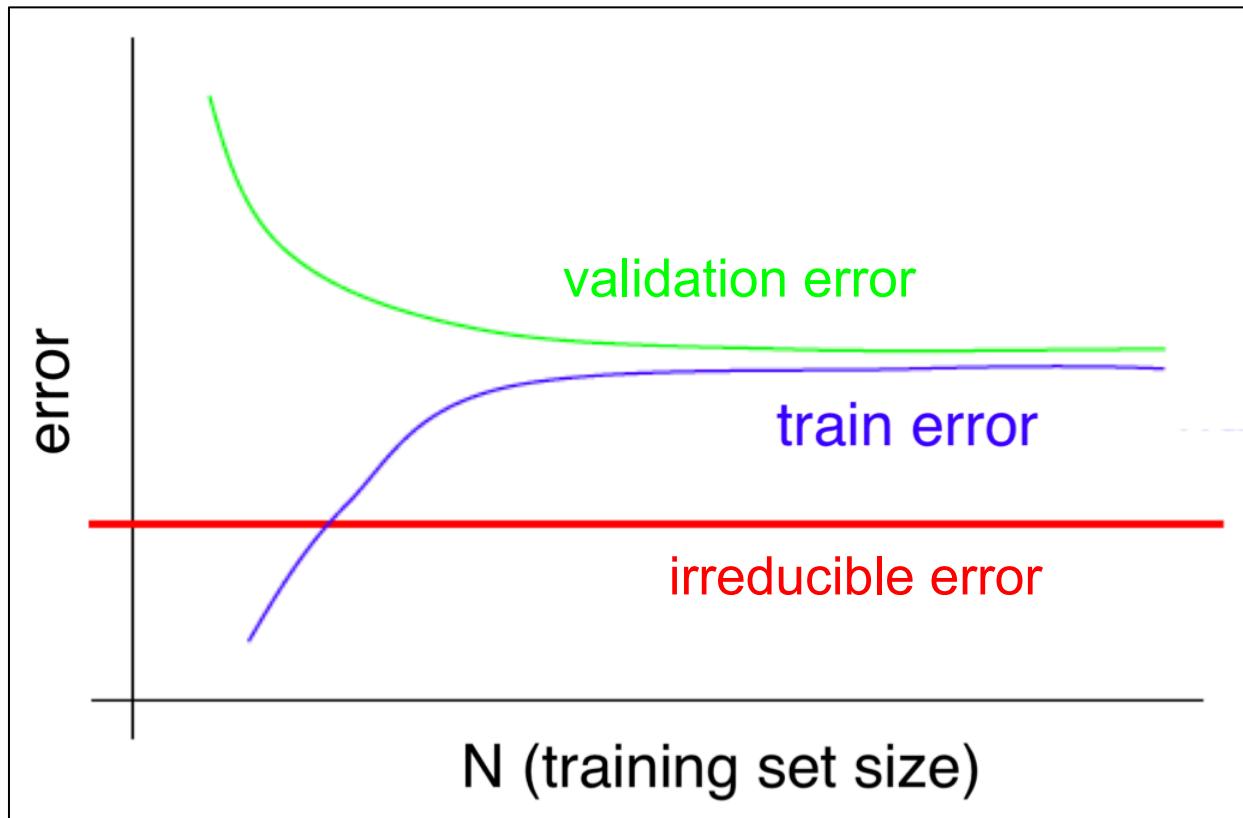


Complexity

## Analysis with Irreducible Error

- The theoretically achievable performance is often unknown.
- Except when a human can solve all samples 100% accurately.
- But even if it is unknown, for tasks where humans are very good, their performance can be used to deduce the irreducible error.
- We now have two characteristic measures which we can read from our learning curves:
  1. Difference between Irreducible error and Training error
  2. Difference between Training error and Validation error

# Learning Curve for Underfitting

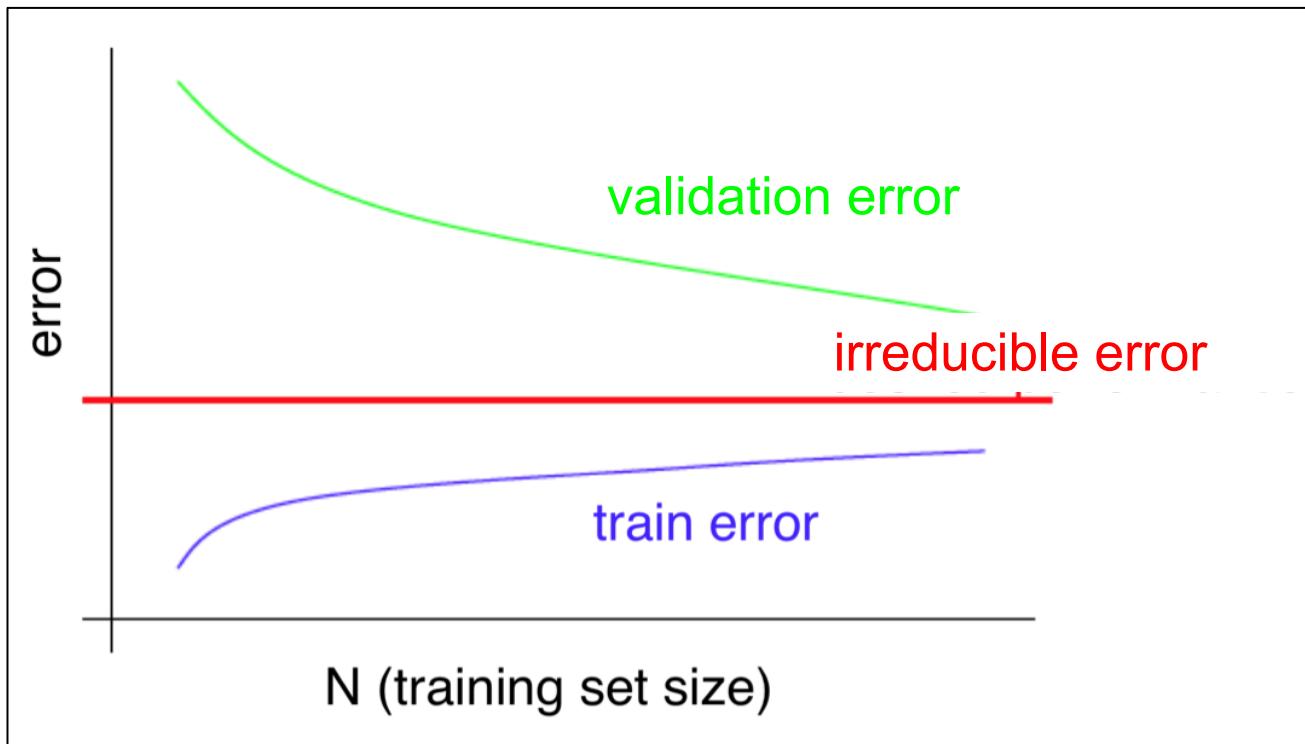


Bad performance on training and test data

# Counteracting Underfitting

- Models that underfit the data are often easier to repair
- If your model misses the relevant structure, try make it more complex
  - Try a larger set of features
  - Try a different set of features
  - Try more complex machine learning algorithms
  - ...

# Learning Curve for Overfitting

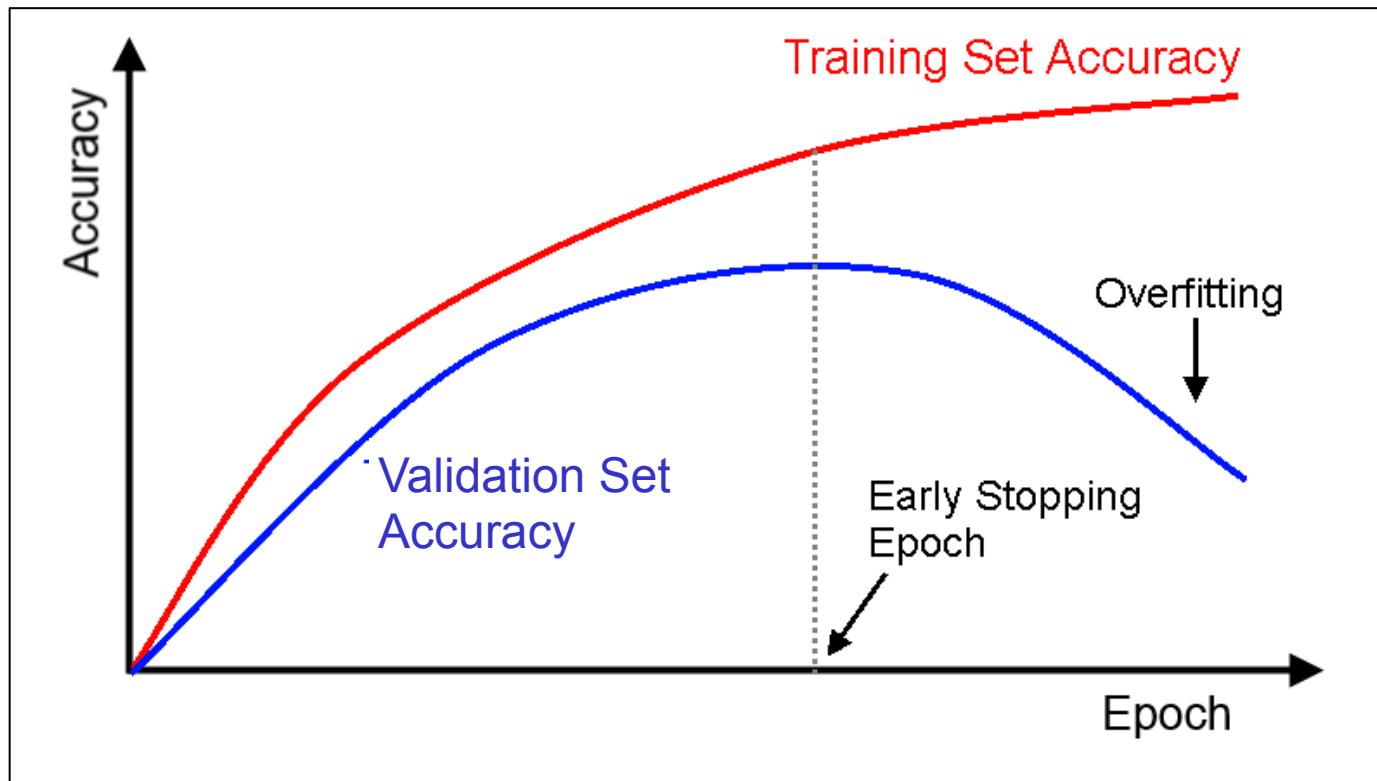


Good performance on training but bad performance on test data

# Counteracting Overfitting

- This happens more frequently and is much harder to repair
- If your model is too sensitive to the training data
  - Try getting more training data
  - Try better data cleaning (e.g. eliminate outliers)
  - Try a smaller set of features through feature selection
  - Try early stopping during iterative training (see next slide)
  - Try tuning the regularization parameter
  - Try ensemble methods
  - ...

# Early Stopping



Especially important for Deep Learning !

# What happened in these two Examples?

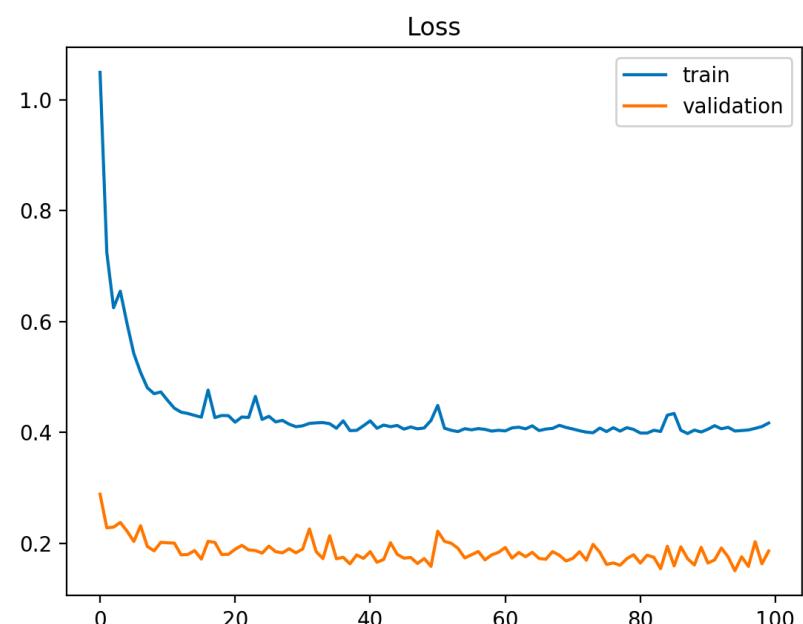
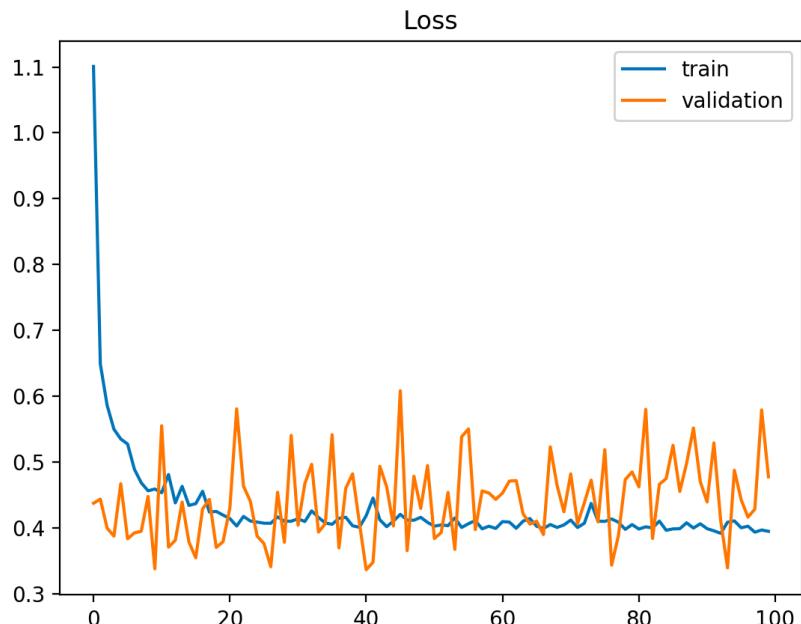


Image source: [machinelearningmastery.com](https://machinelearningmastery.com)

## Analyse Distribution Mismatch

- If it is unclear what part of your error is due to variance and what part due to a distribution mismatch, you can do the following:
- Take your training set only and split it into a training-train and a training-validation set, then build and evaluate your existing model on that.
- You then have eliminated the mismatch (provided your splitting method is correct) and can get a notion of how much your model overfits.

## Example of How to Split

- You want to build a mobile app that recognizes cats
- You have a bunch of low-resolution cat/not cat pictures taken with mobile cameras, but not enough to train your model
- So you scrape the web and download millions of high-resolution cat/notcat images.
- How do you build your training, validation and test sets?
- **Only use mobile camera data:** Does not work, not enough data
- **Split randomly:** Does not work, your score is too optimistic, since your end goal is to classify mobile camera pictures with the app
- **Better:** Exclusively use mobile pictures for validation and test sets, use all high-res pictures and a few mobile pictures for training.

# Error Analysis in Classification

- Pick 100 mislabeled items and analyse them:  
Make a table of the n most frequently occurring problem types (100 x n)
- A problem type might be: Small dogs are mislabeled as cats
- Another one might be: Has problems where head is sideways
- Fix the most frequently occurring problem first
- Mislabeled data can be its own type of problem. If this occurs frequently enough, fix it (but don't forget the lucky guesses).  
It is more important to fix validation and test set labels though.
- Fix only what appears to be a systematic error

# How to Approach a Machine Learning Project

We recommend this approach for industrial projects (not basic research)

1. Data quality assessment is always worth the time
2. Visualize your data
3. Specify and implement a water-proof quality assessment workflow
4. Implement a quick-and-dirty baseline classifier or regressor
5. Run error analysis and diagnostics
  - a) If you have a notion of irreducible error, reduce whichever of the two error measures is bigger first, re-evaluate and repeat until you get close to irreducible error
  - b) If you do not have this notion, reduce bias, then reduce variance, repeat until things do not improve anymore
  - c) If you cannot improve but think there is more potential, perform detailed sample analysis (previous slide)
6. Learn and improve !



## How to get Fired as a Data Scientist – Part 2



In [128]: `data.head()`

Out[128]:

|   | Price  | Year | Cylinders | Gears |
|---|--------|------|-----------|-------|
| 0 | 44800  | 1996 | 8         | 5     |
| 1 | 22800  | 1999 | 8         | 5     |
| 2 | 183710 | 2008 | 8         | 6     |
| 3 | 19900  | 2006 | 4         | 6     |
| 4 | 18999  | 2003 | 5         | 5     |

In [129]: `model = KMeans(n_clusters=3).fit(data)`

```
centers = pd.DataFrame(model.cluster_centers_, columns=data.columns)
centers = centers.round(1)
centers[ 'Cluster Size' ] = np.bincount(model.labels_)

centers
```

Out[129]:

|   | Price    | Year   | Cylinders | Gears | Cluster Size |
|---|----------|--------|-----------|-------|--------------|
| 0 | 17376.6  | 2007.9 | 4.3       | 5.6   | 40836        |
| 1 | 55158.1  | 2011.2 | 5.1       | 6.5   | 13126        |
| 2 | 169479.1 | 2010.4 | 8.5       | 6.7   | 1246         |

In [128]: `data.head()`

Out[128]:

|   | Price  | Year | Cylinders | Gears |
|---|--------|------|-----------|-------|
| 0 | 44800  | 1996 | 8         | 5     |
| 1 | 22800  | 1999 | 8         | 5     |
| 2 | 183710 | 2008 | 8         | 6     |
| 3 | 19900  | 2006 | 4         | 6     |
| 4 | 18999  | 2003 | 5         | 5     |

Do not cluster with  
unnormalized data

In [129]: `model = KMeans(n_clusters=3).fit(data)`

`centers = pd.DataFrame(model.cluster_centers_, columns=data.columns)`

`centers = centers.round(1)`

`centers['Cluster Size'] = np.bincount(model.labels_)`

`centers`

Out[129]:

|   | Price    | Year   | Cylinders | Gears | Cluster Size |
|---|----------|--------|-----------|-------|--------------|
| 0 | 17376.6  | 2007.9 | 4.3       | 5.6   | 40836        |
| 1 | 55158.1  | 2011.2 | 5.1       | 6.5   | 13126        |
| 2 | 169479.1 | 2010.4 | 8.5       | 6.7   | 1246         |

In [45]: `data.head()`

Out[45]:

|   | Price  | Year | Cylinders | Gears | Pink  |
|---|--------|------|-----------|-------|-------|
| 0 | 44800  | 1996 | 8         | 5     | False |
| 1 | 22800  | 1999 | 8         | 5     | False |
| 2 | 183710 | 2008 | 8         | 6     | False |
| 3 | 19900  | 2006 | 4         | 6     | False |
| 4 | 18999  | 2003 | 5         | 5     | False |

In [46]: `data['Pink'].value_counts()`

Out[46]:

```
False      55190
True        18
Name: Pink, dtype: int64
```

In [47]:

```
train, test = train_test_split(data, test_size=0.2)
X_train, X_test = train.drop('Pink', axis=1), test.drop('Pink', axis=1)

scaler = StandardScaler().fit(X_train)
X_train[X_train.columns] = scaler.transform(X_train)
X_test[X_test.columns] = scaler.transform(X_test)
```

In [48]:

```
knn = KNeighborsClassifier(n_neighbors=5, weights='uniform')
knn.fit(X_train, train.Pink)
```

```
print('Accuracy %.5f: ready for production !' % accuracy_score(test.Pink, knn.predict(X_test)))
```

Accuracy 0.99946: ready for production !

In [45]: `data.head()`

Out[45]:

|   | Price  | Year | Cylinders | Gears | Pink  |
|---|--------|------|-----------|-------|-------|
| 0 | 44800  | 1996 | 8         | 5     | False |
| 1 | 22800  | 1999 | 8         | 5     | False |
| 2 | 183710 | 2008 | 8         | 6     | False |
| 3 | 19900  | 2006 | 4         | 6     | False |
| 4 | 16999  | 2003 | 5         | 5     | False |

In [46]: `data['Pink'].value_counts()`

Out[46]:

```
False    55190
True     18
Name: Pink, dtype: int64
```

In [47]:

```
train, test = train_test_split(data, test_size=0.2)
X_train, X_test = train.drop('Pink', axis=1), test.drop('Pink', axis=1)

scaler = StandardScaler().fit(X_train)
X_train[X_train.columns] = scaler.transform(X_train)
X_test[X_test.columns] = scaler.transform(X_test)
```

In [48]:

```
knn = KNeighborsClassifier(n_neighbors=5, weights='uniform')
knn.fit(X_train, train.Pink)
```

```
print('Accuracy %.5f: ready for production !' % accuracy_score(test.Pink, knn.predict(X_test)))
```

Accuracy 0.99946: ready for production !

Do not use accuracy in the presence of skewed data

In [404]: `data.head()`

Out[404]:

|   | Year | Cylinders | Gears | Price  |
|---|------|-----------|-------|--------|
| 0 | 1996 | 8         | 5     | 44800  |
| 1 | 1999 | 8         | 5     | 22800  |
| 2 | 2008 | 8         | 6     | 183710 |
| 3 | 2006 | 4         | 6     | 19900  |
| 4 | 2003 | 5         | 5     | 18999  |

```
In [406]: scaler = StandardScaler().fit(data.drop('Price', axis=1))

train, test = train_test_split(data, test_size=0.2)
X_train, X_test = train.drop('Price', axis=1), test.drop('Price', axis=1)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)

model = KNeighborsRegressor(n_neighbors=1)
model.fit(X_train, train.Price)
print('R2 is %.2f' % r2_score(test.Price, model.predict(X_test)))
```

R2 is 0.59

In [404]: `data.head()`

Out[404]:

|   | Year | Cylinders | Gears | Price  |
|---|------|-----------|-------|--------|
| 0 | 1996 | 8         | 5     | 44800  |
| 1 | 1999 | 8         | 5     | 22800  |
| 2 | 2008 | 8         | 6     | 183710 |
| 3 | 2006 | 4         | 6     | 19900  |
| 4 | 2003 | 5         | 5     | 18999  |

Test set used to determine  
normalization parameters

In [406]: `scaler = StandardScaler().fit(data.drop('Price', axis=1))`

`train, test = train_test_split(data, test_size=0.2)`

`X_train, X_test = train.drop('Price', axis=1), test.drop('Price', axis=1)`

`X_train = scaler.transform(X_train)`

`X_test = scaler.transform(X_test)`

`model = KNeighborsRegressor(n_neighbors=1)`

`model.fit(X_train, train.Price)`

`print('R2 is %.2f' % r2_score(test.Price, model.predict(X_test)))`

R2 is 0.59

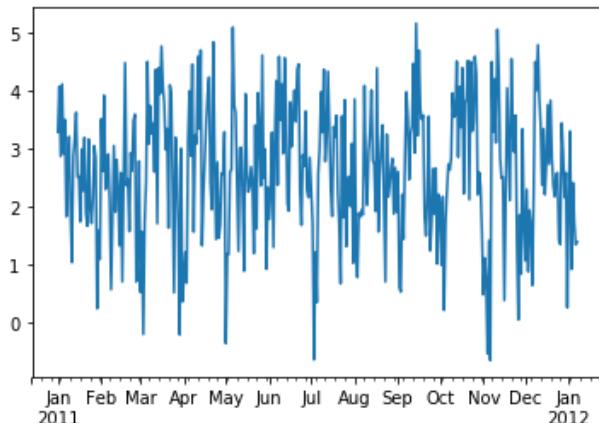
```
In [433]: data.head()
```

```
Out[433]:
```

|            | state | cycle | target   |
|------------|-------|-------|----------|
| 2011-01-01 | 5     | -1.14 | 3.294526 |
| 2011-01-02 | 5     | 0.05  | 4.084435 |
| 2011-01-03 | 8     | 0.50  | 2.872311 |
| 2011-01-04 | 7     | 1.09  | 4.122563 |
| 2011-01-05 | 9     | 0.05  | 2.917532 |

```
In [435]: ts.target.plot()
```

```
Out[435]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16b43d30>
```



```
In [445]: train, test = train_test_split(ts)
X_train, X_test = train.drop('target', axis=1), test.drop('target', axis=1)
model = LinearRegression().fit(X_train, train.target)
print('R2 = %.2f is terribly bad - so is the conceptual error in this code' % r2_score(test.target, model.predict(X_test)))
```

```
R2 = -0.05 is terribly bad - so is the conceptual error in this code
```

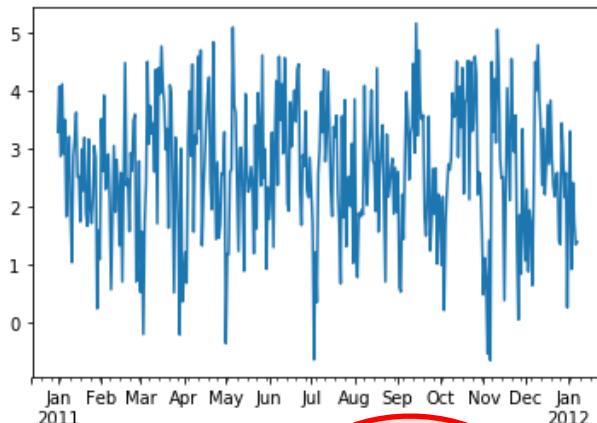
```
In [433]: data.head()
```

```
Out[433]:
```

|            | state | cycle | target   |
|------------|-------|-------|----------|
| 2011-01-01 | 5     | -1.14 | 3.294526 |
| 2011-01-02 | 5     | 0.05  | 4.084435 |
| 2011-01-03 | 8     | 0.50  | 2.872311 |
| 2011-01-04 | 7     | 1.09  | 4.122563 |
| 2011-01-05 | 9     | 0.05  | 2.917532 |

```
In [435]: ts.target.plot()
```

```
Out[435]: <matplotlib.axes._subplots.AxesSubplot at 0x1a16b43d30>
```



```
In [445]: train, test = train_test_split(ts)
X_train, X_test = train.drop('target', axis=1), test.drop('target', axis=1)
model = LinearRegression().fit(X_train, train.target)
print('R2 = %.2f is terribly bad - so is the conceptual error in this code' % r2_score(test.target, model.predict(X_test)))
```

R2 = -0.05 is terribly bad - so is the conceptual error in this code

## Randomized splitting in case of time series

```
df.head()
```

|   | dna_hash | visit_number | blood_pressure | cholesterol | medication | outcome |
|---|----------|--------------|----------------|-------------|------------|---------|
| 0 | 14251    | 1            | 120            | 200         | 1          | 0       |
| 1 | 14251    | 2            | 122            | 205         | 1          | 0       |
| 2 | 23416    | 1            | 130            | 250         | 2          | 1       |
| 3 | 23416    | 2            | 128            | 240         | 2          | 1       |
| 4 | 32525    | 1            | 110            | 180         | 3          | 1       |

```
X_train, X_test, y_train, y_test = train_test_split(df.drop("outcome", axis=1),  
                                                df.outcome)
```

```
clf = RandomForestClassifier()  
clf.fit(X_train, y_train)  
accuracy_score(y_test, clf.predict(X_test))
```

```
df.head()
```

|   | dna_hash | visit_number | blood_pressure | cholesterol | medication | outcome |
|---|----------|--------------|----------------|-------------|------------|---------|
| 0 | 14251    | 1            | 120            | 200         | 1          | 0       |
| 1 | 14251    | 2            | 122            | 205         | 1          | 0       |
| 2 | 23416    | 1            | 130            | 250         | 2          | 1       |
| 3 | 23416    | 2            | 128            | 240         | 2          | 1       |
| 4 | 32525    | 1            | 110            | 180         | 3          | 1       |

```
X_train, X_test, y_train, y_test = train_test_split(df.drop("outcome", axis=1),  
                                                    df.outcome)  
  
clf = RandomForestClassifier()  
clf.fit(X_train, y_train)  
accuracy_score(y_test, clf.predict(X_test))
```

In the presence of grouped data, use GroupKfold or  
GroupShuffleSplit

In [23]: `data.head()`

Out[23]:

|   | Year | Cylinders | Gears | Price  |
|---|------|-----------|-------|--------|
| 0 | 1996 | 8         | 5     | 44800  |
| 1 | 1999 | 8         | 5     | 22800  |
| 2 | 2008 | 8         | 6     | 183710 |
| 3 | 2006 | 4         | 6     | 19900  |
| 4 | 2003 | 5         | 5     | 18999  |

In [32]: `model = [SVR(), LinearRegression(), DecisionTreeRegressor()]  
names = ["Support Vector Machine", "Linear Regression", "Decision Tree"]`

```
for index in range(0,3) :  
  
    train, test = train_test_split(data, test_size=0.2)  
    X_train, X_test = train.drop('Price', axis=1), test.drop('Price', axis=1)  
  
    scaler = StandardScaler().fit(X_train)  
    X_train, X_test = scaler.transform(X_train), scaler.transform(X_test)  
  
    model[index].fit(X_train, train.Price)  
    print('R2 of ' + names[index] + ' is %.2f' % r2_score(test.Price, model[index].predict(X_test)))
```

R2 of Support Vector Machine is -0.14

R2 of Linear Regression is 0.58

R2 of Decision Tree is -0.02

Conclusion: Linear regression performs best on the test set

In [23]: `data.head()`

Out[23]:

|   | Year | Cylinders | Gears | Price  |
|---|------|-----------|-------|--------|
| 0 | 1996 | 8         | 5     | 44800  |
| 1 | 1999 | 8         | 5     | 22800  |
| 2 | 2008 | 8         | 6     | 183710 |
| 3 | 2006 | 4         | 6     | 19900  |
| 4 | 2003 | 5         | 5     | 18999  |

Split inside loops makes  
evaluation incomparable

In [32]: `model = [SVR(), LinearRegression(), DecisionTreeRegressor()]  
names = ["Support Vector Machine", "Linear Regression", "Decision Tree"]`

```
for index in range(0,3) :  
  
    train, test = train_test_split(data, test_size=0.2)  
    X_train, X_test = train.drop('Price', axis=1), test.drop('Price', axis=1)  
  
    scaler = StandardScaler().fit(X_train)  
    X_train, X_test = scaler.transform(X_train), scaler.transform(X_test)  
  
    model[index].fit(X_train, train.Price)  
    print('R2 of '+names[index]+' is %.2f' % r2_score(test.Price, model[index].predict(X_test)))
```

R2 of Support Vector Machine is -0.14

R2 of Linear Regression is 0.58

R2 of Decision Tree is -0.02

Conclusion: Linear regression performs best on the test set

```
# create dataset (no bug here)
(x_train, y_train), (x_val, y_val) = keras.datasets.mnist.load_data()
x_train = x_train.reshape(60000, 784).astype("float32") / 255
x_test = x_test.reshape(10000, 784).astype("float32") / 255
y_train = y_train.astype("float32")
y_test = y_test.astype("float32")
```

```
# create model (no bugs in model architecture)
inputs = keras.Input(shape=(784,), name="digits")
x = layers.Dense(64, activation="relu", name="dense_1")(inputs)
x = layers.Dense(64, activation="relu", name="dense_2")(x)
outputs = layers.Dense(10, activation="softmax", name="predictions")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

```
# compile the model (no bugs here)
model.compile(optimizer="rmsprop", loss="sparse_categorical_crossentropy",
              metrics=["sparse_categorical_accuracy"])
```

```
# run the cross validation (no bug in any of the parameters)
kfold = KFold(n_splits=10, shuffle=True)
scores = []
for train_idx, test_idx in kfold.split(x_train, y_train):
    model.fit(x_train[train_idx], y_train[train_idx], batch_size=64, epochs=1)
    scores.append(model.evaluate(x_train[test_idx], y_train[test_idx], batch_size=128))
```

```
# create dataset (no bug here)
(x_train, y_train), (x_val, y_val) = keras.datasets.mnist.load_data()
x_train = x_train.reshape(60000, 784).astype("float32") / 255
x_test = x_test.reshape(10000, 784).astype("float32") / 255
y_train = y_train.astype("float32")
y_test = y_test.astype("float32")
```

```
# create model (no bugs in model architecture)
inputs = keras.Input(shape=(784,), name="digits")
x = layers.Dense(64, activation="relu", name="dense_1")(inputs)
x = layers.Dense(64, activation="relu", name="dense_2")(x)
outputs = layers.Dense(10, activation="softmax", name="predictions")(x)
model = keras.Model(inputs=inputs, outputs=outputs)
```

```
# compile the model (no bugs here)
model.compile(optimizer="rmsprop", loss="sparse_categorical_crossentropy",
               metrics=[ "sparse_categorical_accuracy"])
```

Model not  
properly

reset, each  
subsequent

fold

continues  
training

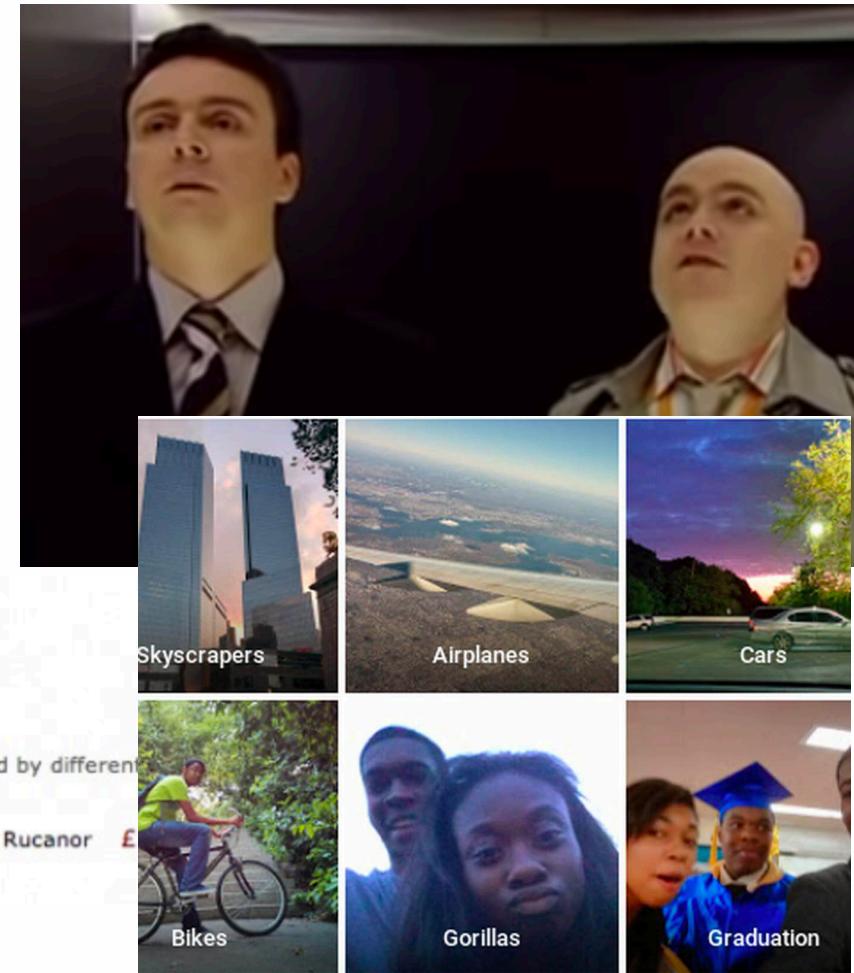
```
# run the cross validation (no bug in any of the parameters)
kfold = KFold(n_splits=10, shuffle=True)
scores = []
for train_idx, test_idx in kfold.split(x_train, y_train):
    model.fit(x_train[train_idx], y_train[train_idx], batch_size=64, epochs=1)
    scores.append(model.evaluate(x_train[test_idx], y_train[test_idx], batch_size=128))
```

# What can go wrong after deployment?

AI Fails will happen!



@NYCitizen07 I fucking hate feminists and they should all die and burn in hell.



## Frequently Bought Together



Price For Both: £20.07

Add both to Basket

These items are dispatched from and sold by different sellers

- This item: Rucanor Aluminium Baseball Bat, Silver - 60 cm Rucanor £12.99
- Balaclava 3 hole black MFH £2.08

# Good Luck!



# Questions ?

